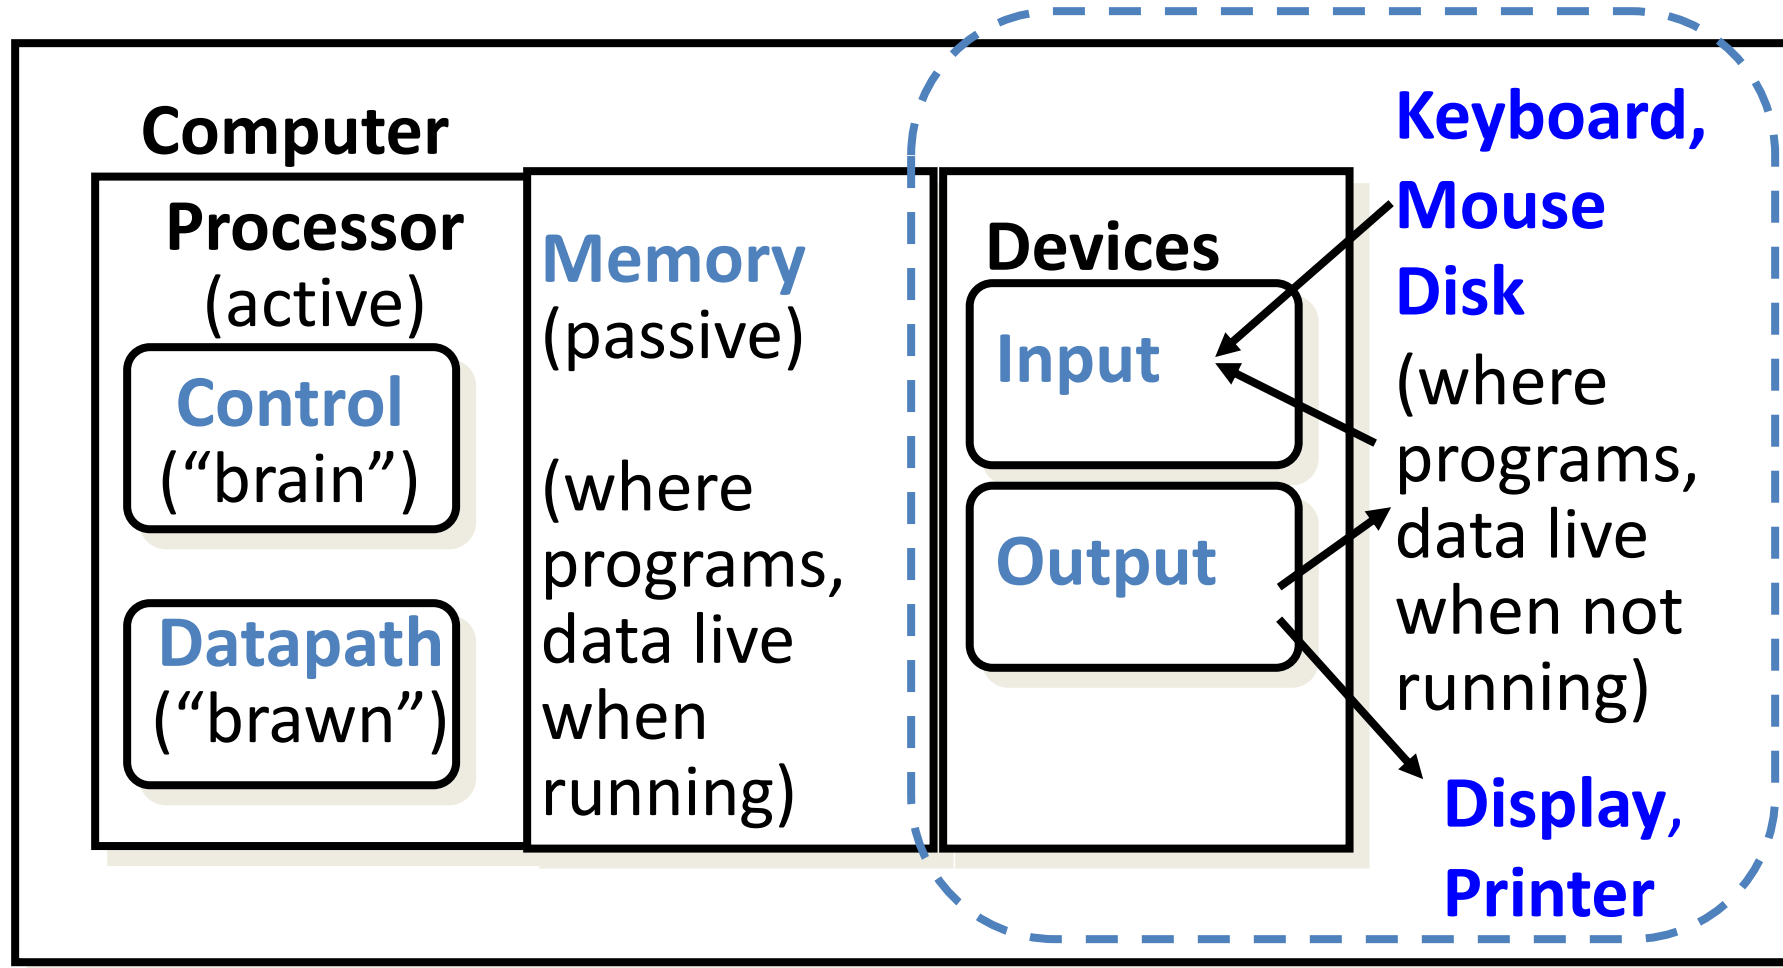


Input/Output: Polling and Interrupts

Outline

- I/O Background
- Polling
- Interrupts
- DMI/DMA

Anatomy: 5 components of any Computer



Motivation for Input/Output

- I/O is how humans interact with computers
- I/O gives computers long-term memory
- I/O lets computers do amazing things



http://www.youtube.com/watch?v=w3n4D_VyXN8


I/O Device Examples and Speeds

- Kilobytes transferred per second from mouse to display... million to one range of rates!

| Device | Behaviour | Partner | Data Rate KB/s |
|------------------|--------------|---------|----------------|
| Keyboard | Input | Human | 0.01 |
| Mouse | Input | Human | 0.02 |
| Voice output | Output | Human | 5.00 |
| Floppy disk | Storage | Machine | 50.00 |
| Laser Printer | Output | Human | 100.00 |
| Magnetic Disk | Storage | Machine | 10,000.00 |
| Network-LAN | Input/Output | Machine | 10,000.00 |
| Graphics Display | Output | Human | 30,000.00 |

What do we need to make I/O work?

- A way to connect many types of devices to the Proc-Mem
- A way to control these devices, respond to them, and transfer data
- A way to present them to user programs so they are useful



Processor-I/O Speed Mismatch


- What must the processor do for I/O?
 - Input: reads a sequence of bytes
 - Output: writes a sequence of bytes
- 1 GHz microprocessor can execute 1 billion load or store instructions per second.
 - I/O devices data rates range from 0.01 KB/s to 30 MB/s
- Input: device may not be ready to send data as fast as the processor loads it
 - Also, might be waiting for human to act
- Output: device not be ready to accept data as fast as processor stores it
- What to do?

Paradigms

- There are three common paradigms for interfacing with a I/O devices:
 - Polling
 - Sending bytes one-at-a-time.
 - Cheapest solutions and easiest to build
 - Interrupts
 - Ignoring the device until it is ready.
 - Direct Memory Interfacing/Accessing (DMI/DMA)
 - Almost automatic & runs in parallel
 - Most expensive solutions, easy to program, but hard to build

Polling Paradigm

- The controller is attached to RAM. This is called **Memory Mapped I/O**
- Certain addresses are not regular memory
- Instead, these addresses correspond to registers in I/O devices – **LW/SW**
- Used by MIPS.





`sw $t1, ADDRESS`

Memory Mapped

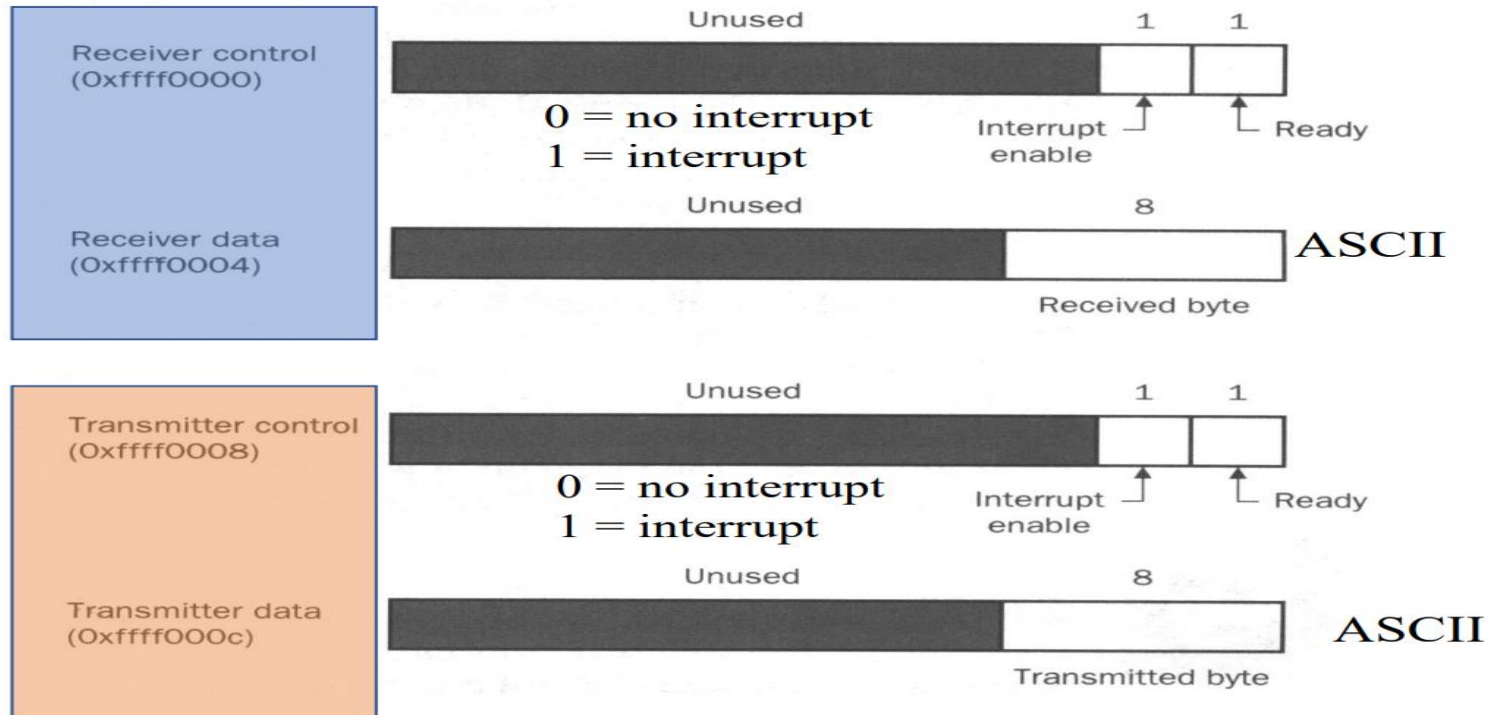
A portion of RAM's memory is physically connected to the address space of the peripheral registers (on-board and off-board ~ via slots)

Processor Checks Status before Acting

- Path to device generally has 2 registers:
 - [Control Register](#), says it's OK to read/write (I/O ready)
 - [Data Register](#), contains data
- Processor reads from Control Register in loop, waiting for device to set *Ready* bit in Control register (0 becomes 1) to say its OK
- Processor then loads from (input) or writes to (output) data register
 - Load from or Store into Data Register resets Ready bit (1 becomes 0) of Control Register

I/O Simulation

- MARS simulate one I/O device:
 - Memory-mapped terminal (keyboard + display)
 - Read from keyboard ([receiver](#)); 2 device registers
 - Writes to terminal ([transmitter](#)); 2 device registers



I/O Control and Data Registers

- Control register rightmost bit (0): Ready
 - Receiver: Ready==1 means character in Data Register has not yet been read; the 1 changes to 0 when data is read from Data Register
 - Transmitter: Ready==1 means transmitter is ready to accept a new character; 0 means the transmitter is still busy writing last char
- The I.E. (*interrupt enable*) bit we'll discuss later
- Data register rightmost byte has data
 - Receiver: last char from keyboard; other bytes in word are zero
 - Transmitter: when we write the rightmost byte, this will write the char to the display

I/O Example

- Input: Read from keyboard into \$v0 (**#GETCHAR#**)

```
Waitloop:      lui    $t0, 0xffff          # memory address 0xffff0000
               lw     $t1, 0($t0)         # receiver control
               andi   $t1,$t1,0x0001      # check ready bit with mask
               beq    $t1,$zero, Waitloop
               lw     $v0, 4($t0)         # data
```

- Output: Write to display from \$a0 (**#PUTCHAR#**)

```
Waitloop:      lui    $t0, 0xffff          # memory address 0xffff0000
               lw     $t1, 8($t0)         # transmitter control
               andi   $t1,$t1,0x0001      # check ready bit with mask
               beq    $t1,$zero, Waitloop
               sw     $a0, 12($t0)        # data
```

- Processor waiting for I/O called *Polling*

Cost of Polling? - mouse

- Assume a processor with a 1 GHz clock takes 400 clock cycles for a polling operation (call polling routine, accessing the device, and returning)
- Determine % of processor time for polling...
 - Mouse polled 30 Hz so as not to miss user movement



% Processor time to poll mouse

- Mouse Polling, Clocks/sec

$$= 30 \times 400$$

$$= 12000 \text{ clocks/sec}$$

- % Processor for polling:

$$\frac{12 \times 10^3}{1 \times 10^9} = 0.0012\%$$

- Polling mouse little impact on processor

Cost of Polling? – Floppy disk

- Assume a processor with a 1 GHz clock takes 400 clock cycles for a polling operation (call polling routine, accessing the device, and returning)
- Determine % of processor time for polling...
 - Floppy disk transfers data in 2-Byte units
 - Has a data rate of 50 KB/second
 - No data transfer can be missed



% Processor time to poll floppy disk

- Frequency of Polling Floppy
= 50 KB/s / 2B = 25K polls/sec
- Floppy Polling, Clocks/sec
= 25K × 400 = 10,000,000 clocks/sec
- % Processor for polling:
$$\frac{10 \times 10^6}{1 \times 10^9} = 1\%$$
- OK if not too many I/O devices

Cost of Polling? – Hard disk

- Assume a processor with a 1 GHz clock takes 400 clock cycles for a polling operation (call polling routine, accessing the device, and returning)
- Determine % of processor time for polling...
 - Hard disk transfers data in 16-Byte chunks
 - Can transfer at 16 MB/second
 - Again, no transfer can be missed



% Processor time to hard disk

- Frequency of Polling Disk
= 16 MB/s / 16B = 1M polls/sec
- Disk Polling, Clocks/sec
= 1M × 400 = 400,000,000 clocks/sec
- % Processor for polling:
$$\frac{400 \times 10^6}{1 \times 10^9} = 40\%$$
- Unacceptable!

What is the alternative to polling?

- Wasteful to have processor spend most of its time “spin-waiting” for I/O to be ready
- Want an unplanned procedure call that would be invoked only when I/O device is ready
- Solution: use *exception mechanism* to help I/O by *Interrupting* the program when I/O ready, return when done with data transfer


I/O Interrupt

- An I/O interrupt is like overflow exceptions except that
 - An I/O interrupt is “asynchronous”
 - More information needs to be conveyed
- An I/O interrupt is asynchronous with respect to instruction execution
 - It is not associated with any instruction, but it can happen in the middle of any given instruction
 - It does not prevent any instruction from completion

Definitions for Clarification

Exception: signal marking that something “out of the ordinary” has happened and needs to be handled

- Interrupt:
 - asynchronous exception
 - e.g., I/O device.
- Trap:
 - synchronous exception
 - e.g., segmentation faults.
 - Try..catch is an example of language support.

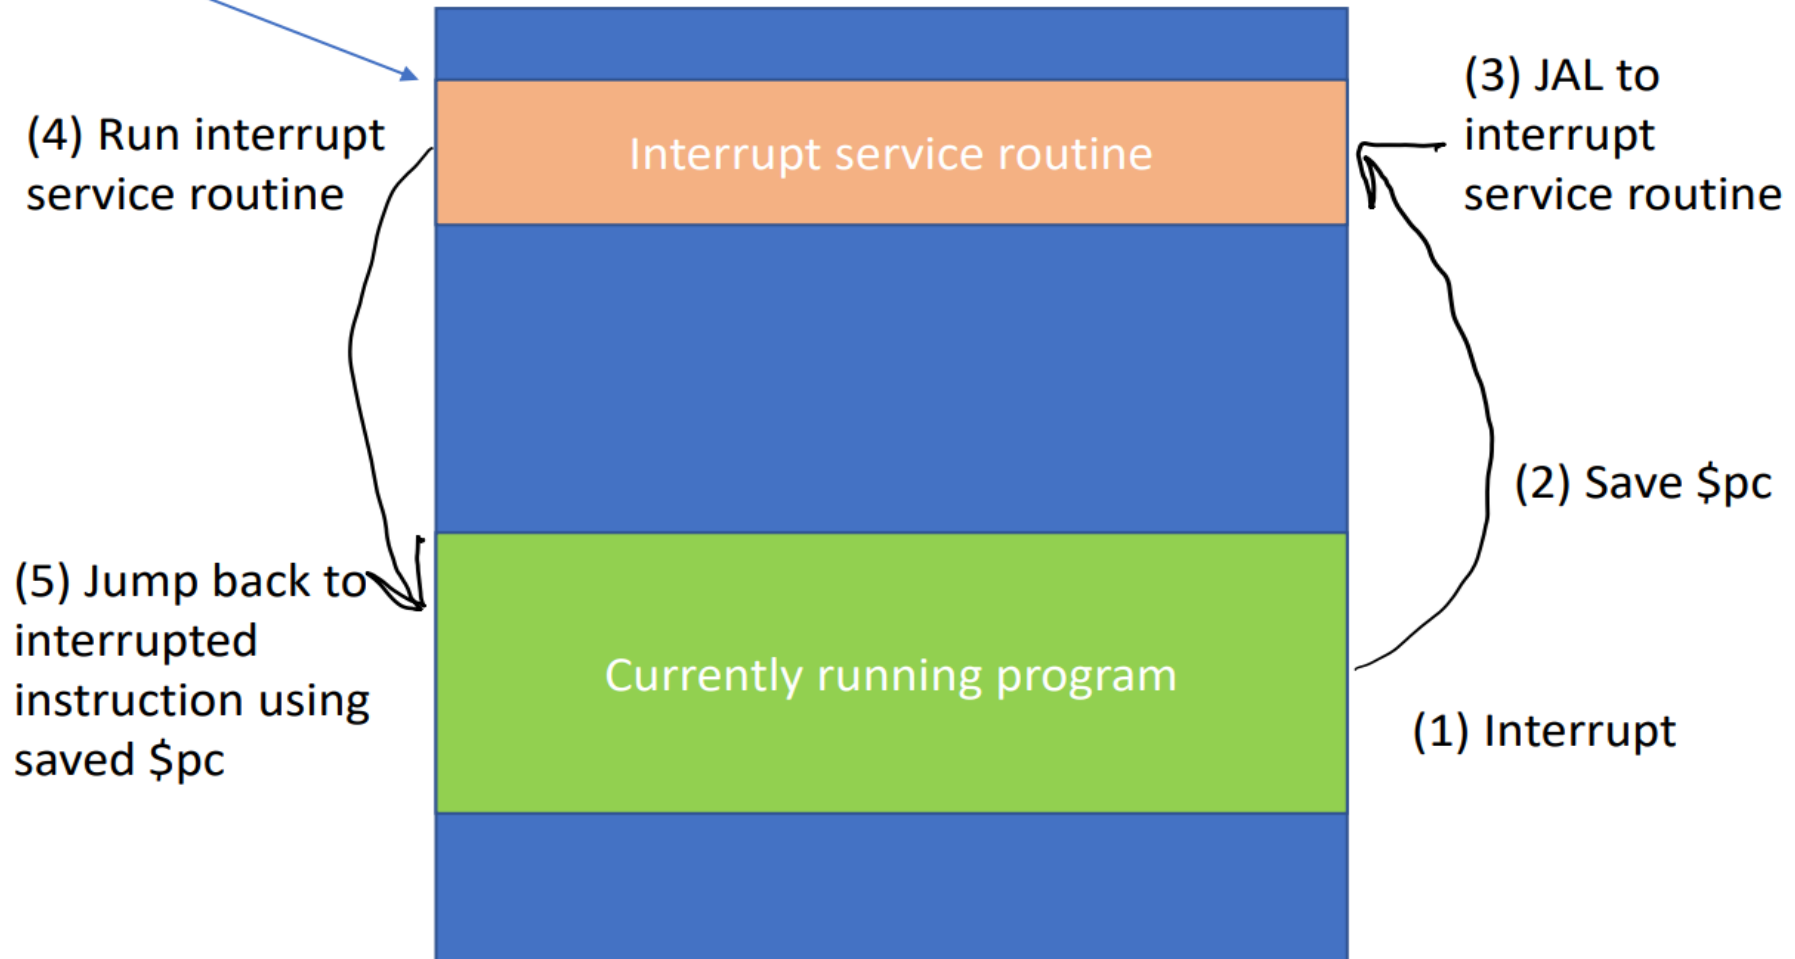


Exception Types

- External events
 - Peripheral wants CPU time
- Memory translation
 - Non-existent memory location
- Run-time errors
 - Divide by zero – Overflow, etc.
- Coding Errors
 - Non-existent op-code

Interrupt Driven Data Transfer

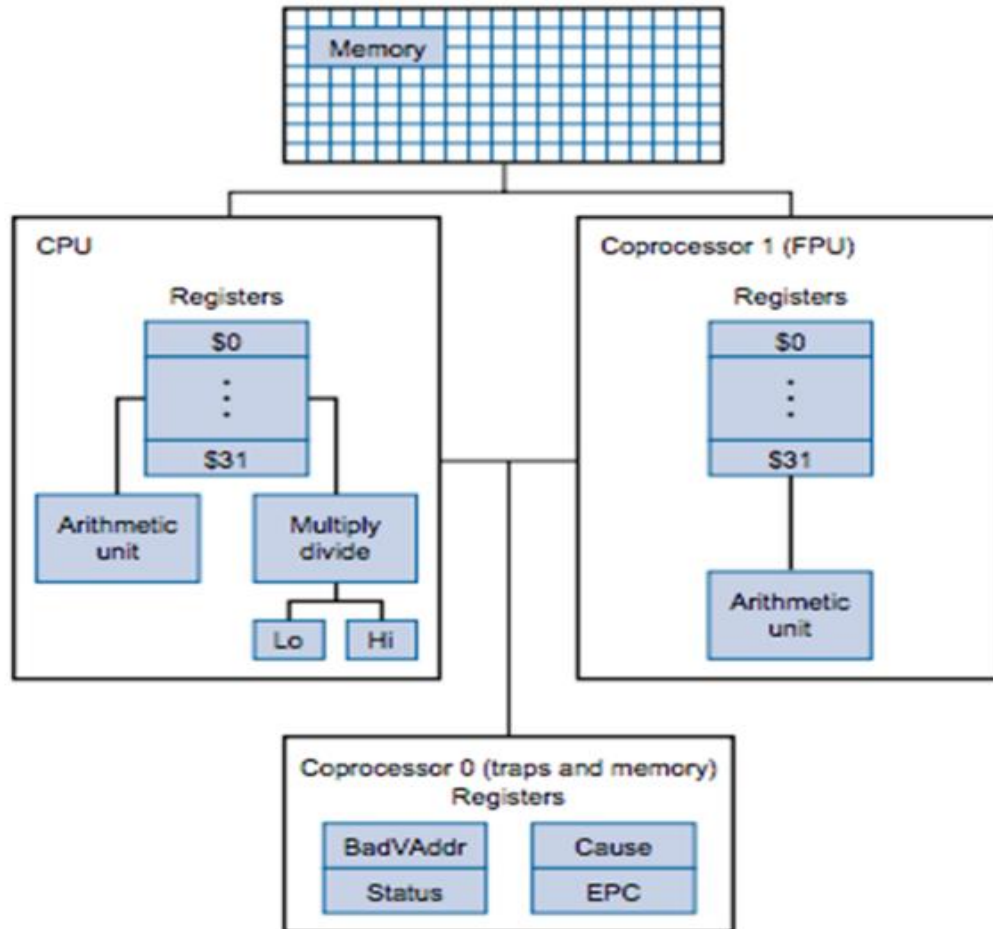
0x80000080 or 0x80000180



Instruction Set Support for I/O Interrupt

- Save the PC for return, but where?
- Where to go when interrupt occurs?
- Determine cause of interrupt?

The CPU Chip Set



- A **chip set** is a group of machines that work together.
- The MIPS CPU has 3 machines: Integer CPU, FPU (Co-1), and Error (Co-0).
- Co-0 and Co-1 have their own registers.

Instruction Set Support for I/O Interrupt

- Portion of MIPS architecture for interrupts called ***coprocessor 0***
- Coprocessor 0 Instructions
 - Data transfer: lwc0 (load word to c0), swc0 (store word from c0)
 - Move: mfc0 (move from c0 to CPU), mtc0 (move CPU to c0)
- Coprocessor 0 Registers:

| Name | Number | Usage |
|--------|--------|------------------|
| Status | \$12 | Interrupt enable |
| Cause | \$13 | Exception type |
| EPC | \$14 | Return address |

Example:

mfc0 \$t1, \$13 # move Cause to t1

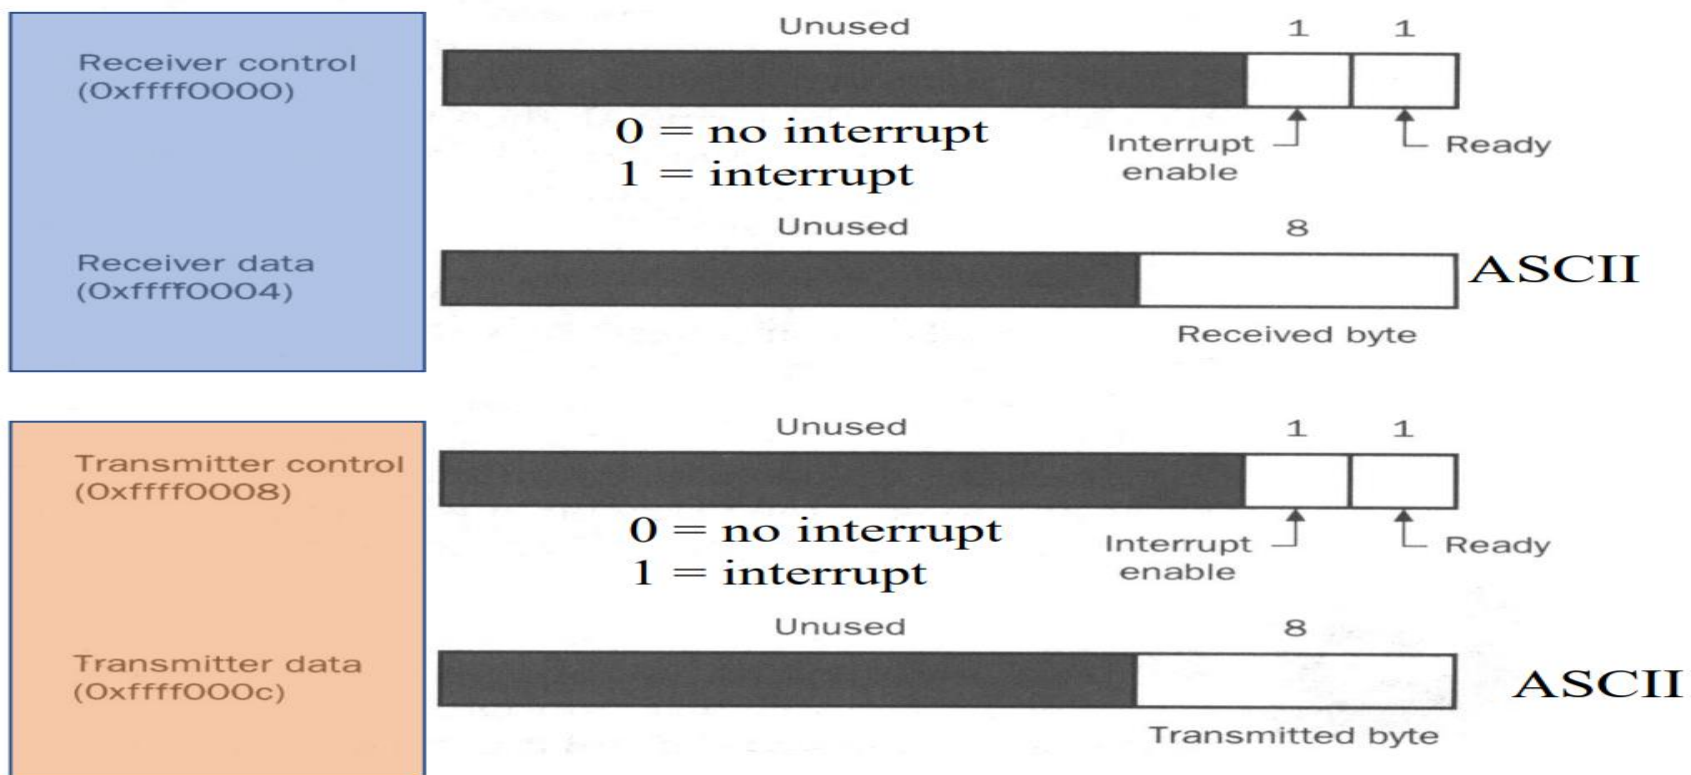
Instruction Set Support for I/O Interrupt

Answers to question in slide#26

- Save the PC for return
 - But where?
 - The address of the instruction that is about to be executed is saved into EPC
- Where to go when interrupt occurs?
 - MIPS defines location: 0x80000180.
 - PC is set to be 0x80000180, the starting address of the interrupt service routine (interrupt handler), which takes the processor to the interrupt handler for execution.
 - The last instruction of the interrupt service routine sets the value of the PC to the value stored in EPC
- Determine cause of interrupt?
 - MIPS has Cause Register, 5-bit field (bits 6 to 2) gives cause of exception

Interrupt Driven I/O Simulation

- I.E. is the **Interrupt Enable** bit... set it to 1 to have interrupt occur whenever Ready bit is set



Benefit of Interrupt-Driven I/O

- Find the % of 1 GHz processor consumed if the hard disk is only active 5% of the time
 - Assuming 500 clock cycle overhead for each transfer, including the interrupt
 - Also assume the hard disk can transfer at 16 MB/s in 16 byte chunks



Benefit of Interrupt-Driven I/O

- Answer:
 - Disk Interrupts/sec = $16 \text{ MB/s} / 16\text{B}$
= 1M interrupts/sec
 - Disk Interrupts, Clocks/sec = $1\text{M} \times 500$
= 500,000,000 clocks/sec
 - % Processor needed during transfer:
$$\frac{500 \times 10^6}{1 \times 10^9} = 50\%$$
- Disk active 5% \Rightarrow 5% \times 50% \Rightarrow 2.5% busy

Questions Raised about Interrupts

- Which I/O device caused exception?
 - Needs to convey the identity of the device generating the interrupt
- Can avoid interrupts during the interrupt routine?
 - What if more important interrupt occurs while servicing this interrupt?
 - Allow interrupt routine to be entered again?
- Who keeps track of status of all the devices, handle errors, know where to put/supply the I/O data?

4 Functions OS must provide

- Provides abstractions for accessing devices by supplying routines that handle low-level device operations
- Handles the exceptions generated by I/O devices (and arithmetic exceptions generated by a program)
- Tries to provide equitable access to the shared I/O resources, as well as schedule accesses to enhance system performance.
- Guarantees that user's program accesses only the portions of I/O device to which user has rights (e.g., file access)

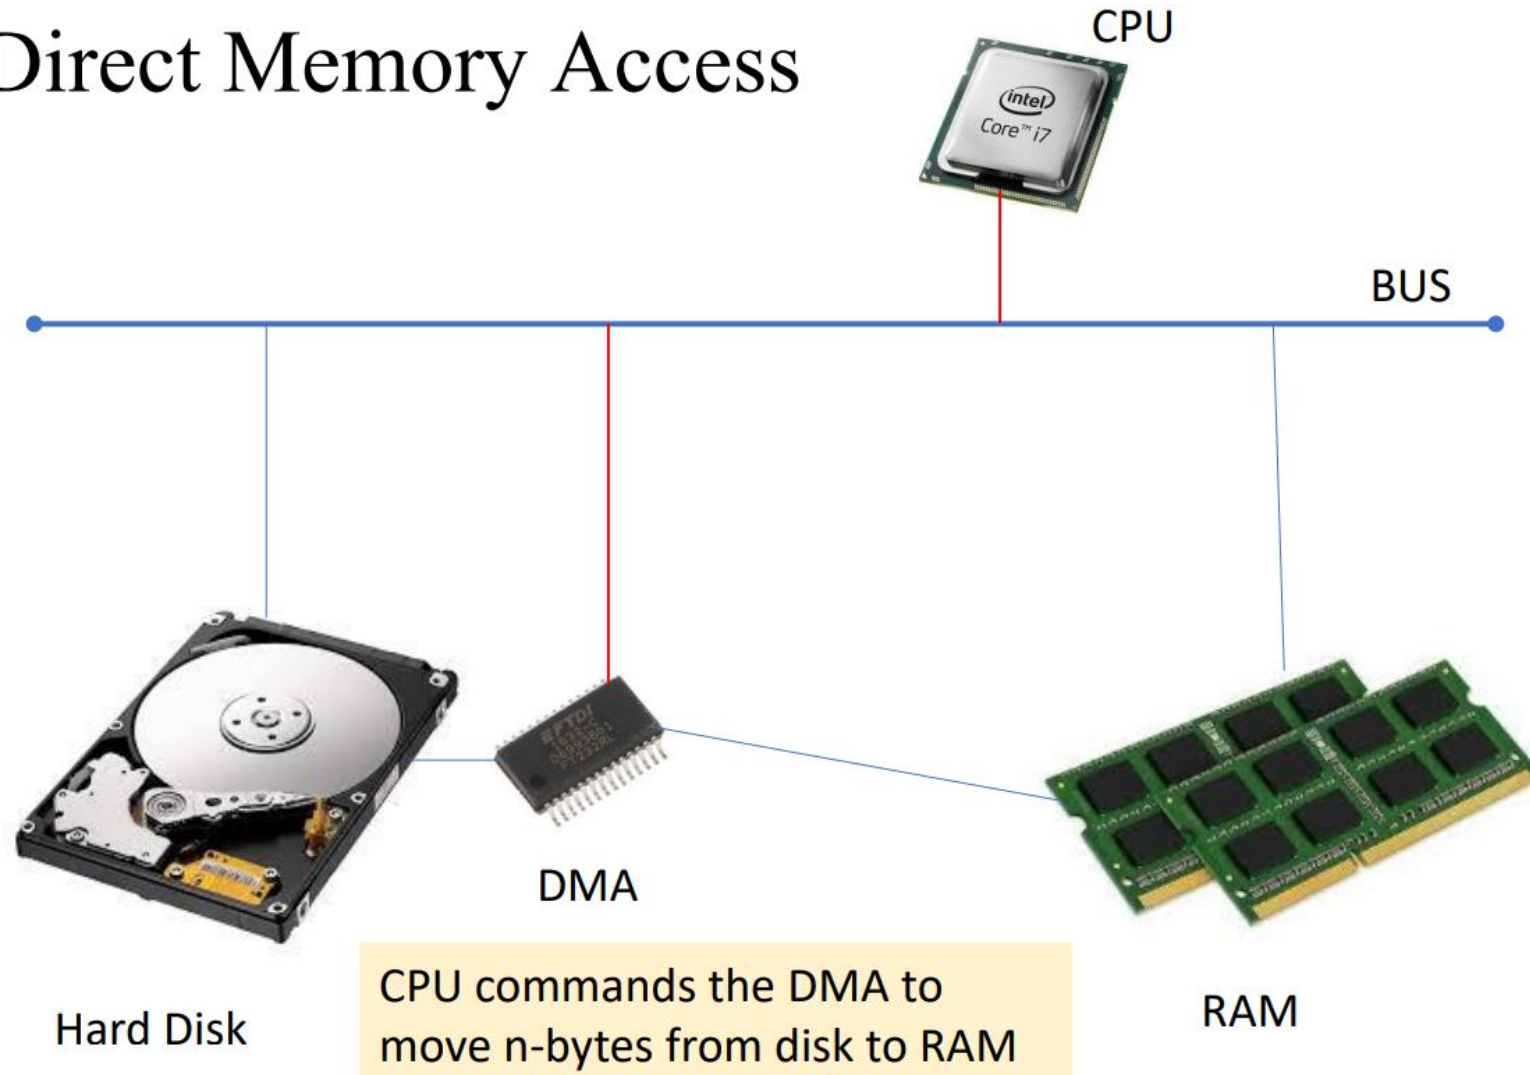
DMI/DMA

What is Direct Memory Interface/Access?

- When I/O device has its own **private** bus and clock that connects it to the RAM.
- This means the CPU does not need to be involved in the copying of data from or to the I/O device.
- Procedure:
 - First, CPU tells DMI device what to do
 - Second, CPU runs programs in parallel to DMI. DMI transfer data to destination when things are available. DMI worries about things.
 - Third, DMI sends interrupt to CPU to say “done”.

DMI/DMA

Direct Memory Access



Benefit of DMI/DMA

- Assume 1GHz CPU, 16MB/sec HDD
- DMA overhead (initialize) 1000 ticks
- DMA Interrupt overhead (process at end) 500 ticks
- DMA transfer rate of 16 bytes per tick.
- Processor usage?

Answer:

DMA transfer = 16MB / 16 bytes/sec = 1M transfers/sec

DMA work = (1000 cs + 500 cs) = 1500 cs

Processor usage = 1500 cs / 1000000000 cs = 0.0000015 = 0.00015%

Only for overhead ... no transfer work

Things to Remember

- I/O gives computers their 5 senses
- I/O speed range is million to one
- Because of the speed of the processor, it must synchronize with I/O devices before using them (reading or writing)
- Polling works, but expensive
 - Processor repeatedly queries devices
- Interrupts work, but are more complex
 - Device causes an exception, causing OS to run to deal with the device
- DMI/DMA work better, but expensive to build.
 - Handing off the data transfer job from and to the I/O device to specialized hardware; that's the DMA/DMI.
- I/O control leads to Operating Systems

Review and More Information

- Textbook 5th edition A7 and A8